

COMMENT [ann_37]: Note: A summary is provided at Pg 12. It is recommended to go through the summary first before reading the full paper.

SECTION [sec_001]: AGENT WORKFLOW MEMORY

AGENT WORKFLOW MEMORY

PARA [para_001]: Zora Zhiruo Wang Jiayuan Mao D...

Zora Zhiruo Wang^C Jiayuan Mao^M Daniel Fried^C Graham Neubig^C
^C Carnegie Mellon University ^M Massachusetts Institute of Technology

SECTION [sec_002]: ABSTRACT

ABSTRACT

Despite the potential of language model-based agents to solve real-world tasks such as web navigation, current methods still struggle with long-horizon tasks with complex action trajectories. In contrast, humans can flexibly solve complex tasks by learning reusable task workflows from past experiences and using them to guide future actions. To build agents that can similarly benefit from this process, we introduce Agent Workflow Memory (AWM), a method for inducing commonly reused routines, i.e., *workflows*, and selectively providing workflows to the agent to guide subsequent generations. AWM flexibly applies to both *offline* and *online* scenarios, where agents induce workflows from training examples beforehand or from test queries on the fly. We experiment on two major web navigation benchmarks — Mind2Web and WebArena — that collectively cover 1000+ tasks from 200+ domains across travel, shopping, and social media, among others. AWM substantially improves the baseline results by 24.6% and 51.1% relative success rate on Mind2Web and WebArena while reducing the number of steps taken to solve WebArena tasks successfully. Furthermore, online AWM robustly generalizes in cross-task, website, and domain evaluations, surpassing baselines from 8.9 to 14.0 absolute points as train-test task distribution gaps widen.

<https://github.com/zorazrw/agent-workflow-memory>

SECTION [sec_003]: 1 INTRODUCTION

1 INTRODUCTION

Language model (LM) based agents are rapidly improving, and are now able to tackle digital tasks such as navigating the web (Zhou et al., 2024; Deng et al., 2023) or operating mobile apps (Rawles et al., 2023, 2024). Current agents mostly integrate a fixed set of given examples via training (Fu et al., 2024; Murty et al., 2024) or in-context learning (Zheng et al., 2024). This allows them to perform well on action sequences similar to those presented in these examples, but they do not learn from past successes and failures, and are therefore unable to adapt over time. Essentially, they fail to grasp the key to disentangling increasingly complex tasks — to extract and learn reusable task workflows shared across similar tasks and environments (Yu et al., 2023; Wang et al., 2024a). Moreover, as agents solve each task separately, they are unable to learn from past successes and failures, and are therefore unable to adapt over time.

Motivated by how humans abstract common task routines from past experiences and apply such knowledge to guide future activities (Chi et al., 1981; 2014), we propose agent workflow memory (AWM) (§2) to realize a similar mechanism in agents. AWM induces workflows from agent trajectories by extracting reusable routines, and then integrates these workflows into agent memory to guide future task-solving processes. Each workflow represents a goal with a common routine extracted from available action trajectories, which allows it to effectively capture the most essential and reusable skills agents need to acquire to successfully solve increasingly complex tasks. As an example, Figure 1 shows workflows induced by AWM on the WebArena map




Figure 1: AWM enables agents to continuously induce and apply workflows to improve performance, compared to stagnant baselines. We show results by AWM on the WebArena map as an example.

SECTION [sec_005]: The equation is: $\sqrt{1} = \sqrt{rac{...}}$

HIGHLIGHT [ann_66]: ark (Zhou et al., 2024). A

test split of the benchmark (Zhou et al., 2024). AWM starts with a basic set of built-in actions and solves new tasks, e.g., learning to “find a place by its name” from a few examples. Moreover, AWM continues to build more complex workflows on top of new experiences and previously acquired workflows. For example, the “find a place by its name” workflow, once induced, effectively serves as a subgoal to build a more complex workflow “get the zip code of a place.” Such continual learning mechanisms create a snowball effect to induce and apply increasingly complex workflows while expanding the agent memory, often yielding a substantial performance gap over a vanilla agent that does not adapt. This gap over the baseline rises as high as 22.5 points on WebArena after rolling over only tens of

COMMENT [ann_73]: What would

COMMENT [ann_74]: module do?

module do?

COMMENT [ann_76]: Online scenarios are

COMMENT [ann_77]: the real test of an

COMMENT [ann_78]: age tic behavio

age tic behaviour

AWM readily operates in both *offline* and *online* scenarios, where annotated examples are either available or non-existent. When high-quality annotated examples are available for a task, AWM operating in an *offline* fashion can extract reusable workflows from these canonical examples and integrate them into memory to assist test-time inference. Even if no examples exist, AWM *online* can also run in a *supervision-free setting*, where it iteratively induces workflows from self-

HIGHLIGHT [ann_69]: We evaluate AWM on two agent web navigation benchmarks by an evaluator comment [ann_79]: §3)

HIGHLIGHT [ann_67]: n-based evaluation (Zhou et al., 2024), and Mind2Web

We evaluate AWM on two agent web navigation benchmarks (§3): WebArena, which provides rigorous execution-based evaluation (Zhou et al., 2024), and Mind2Web, which emphasizes broad tasks and domain coverage (Deng et al., 2023). On WebArena, AWM improves over the top published autonomous method (Drouin et al., 2024) by 51.1% relative success rate, and even outperforms methods augmented with human expert written workflows (Sodhi et al., 2023) by 7.9%. On Mind2Web, AWM effectively improves the cross-task results by 24.6% in relative step-wise success rate.

We further demonstrate the generalizability of AWM on both datasets. On WebArena, we create a cross-template subset where each example is instantiated from different task templates. AWM still consistently surpasses all baseline approaches, demonstrating its reliable cross-task workflow adaptability (§3.1). On Mind2Web, we evaluate AWM on the cross-website and cross-domain test splits to examine its domain generality, where it scores 8.9–14.0 absolute points higher over baseline, and the margins become more substantial as the train-test distribution gap widens (§3.2). Both results show the superior generalization of AWM across tasks, websites, and domains.

§3.1: 2 AGENT WORKFLOW MEMORY

2 AGENT WORKFLOW MEMORY

COMMENT [ann_82]: §2.1)

In this section, we first describe the web navigation task (§2.1), then introduce the workflow representation (§2.2), and describe the mechanism of AWM as well as various instantiations (§2.3).

§2.1: 2.1 PROBLEM STATEMENT

2.1 PROBLEM STATEMENT

HIGHLIGHT [ann_64]: memory M , which contains documents....

HIGHLIGHT [ann_71]: For the purpose of this paper, we consider agents with a language model backbone L and text-based memory M , where the base memory contains documentation of built-in actions such as `CLICK`

and `TYPE`. To solve a task specified by a natural language (NL) instruction q , the agent acts in an environment defined by a transition function T . For each time step t_i , the environment state s_i gives observation o_i , which is then passed into the model to generate action a_i via $L(q, M, o_i) \rightarrow a_i$. The action is executed in the environment and changes the state as $T(s_i, a_i) \rightarrow s_{i+1}$. This observation loop iterates until the model predicts the stop action $a_i = \text{STOP}$, or reaches a task termination condition, e.g., a maximum pre-determined number of steps.

COMMENT [ann_85]: The workflow

COMMENT [ann_87]: memory

COMMENT [ann_88]: execution loop

execution loop

COMMENT [ann_89]: What happens when a task

COMMENT [ann_90]: successfully?

COMMENT [ann_91]: successfully?

successfully?

Each completed task forms an experience e , which comprises an NL instruction q and a trajectory of steps attempting to solve the task, where each step p contains the agent observation o obtained from the current state, and action taken by the agent a , formulated as $p = (o, a)$. Our goal is to induce useful workflows $\mathcal{W} = \{w\}$ from the set of experiences $\mathcal{E} = \{e\}$ constructed from past or collected examples, using an induction module I via $I(\mathcal{E}) \rightarrow \mathcal{W}$. We add induced workflows into the agent memory M as guidance for subsequent task-solving.

COMMENT [ann_92]: How can we le

COMMENT [ann_93]: past experiences?

past experiences?

Next, we introduce the workflow representation design, workflow induction method, and agent memory update with workflows in varied setups.

PARA [para_062]: 1Memory is usually implemented...

1Memory is usually implemented as a system prompt or auxiliary information in the main prompt context.

§3.1: 3.1 The equation is $v = mc^2$ 2...

2.2 WORKFLOW REPRESENTATION

Similar to an experience, a workflow comprises two components: first, a textual description of the workflow d ; and second, a series of steps to finish the workflow P^d .

Workflow Description To present workflows in a format where agents can learn from them properly, it is important to describe the high-level goal of the series of actions. Therefore, we associate each workflow with an NL task description d , essentially a summary of the workflow’s function, by heuristically extracting from experience instructions or summarizing with an LM (see §2.3).

Workflow Trajectory The workflow trajectory contains a series of steps $(p_1, p_2, \dots)$ to finish the process described in d . Each p consists of three parts, demonstrated in Figure 2: shown: (1) A description of the current environment state in NL, such as “Order {id} is shown”; (2) The reasoning process elaborated by the agent to decide which action to generate based on observations, such as “Order {id} is found, I will now terminate the task.”; and (3) an action represented as an executable program over the environment, i.e., `stop()` that realizes termination.

2.3 INDUCING AND USING WORKFLOWS

2.3.1 INDUCING AND USING WORKFLOWS

At the core of AWM is an induction module I that induces a set of workflows $\mathcal{W}$ from one or more past agent experiences $\mathcal{E} = \{e_i\}_{i=1}^m$. Each experience $e = (q, P^e)$ contains an NL task instruction q and an action trajectory P^e that consists of a sequence of steps (observation and action) $P^e = (p_1^e, \dots, p_n^e)$ that were taken to solve q . The workflow induction module operates by taking in $\mathcal{E}$ and producing a set of workflows, as $I(\mathcal{E}) \rightarrow \mathcal{W} = \{w\} = \{(d_j, P_j^d)\}$.

LM-based Workflow Induction To produce workflows that more accurately capture reusable trajectories across tasks, we propose an LM-based module I that prompts the agent to extract common sub-routines from one or more input experiences.

Different from task instructions that specify concrete, less-repetitive tasks, e.g., “Buy dry cat food on Amazon and deliver to my address”, we deliberately prompt models to induce workflows at finer granularities, i.e., a sub-task “search for a product on Amazon” that frequently re-appears as part of multiple similar instructions. Meanwhile, instead of giving example-specific values (e.g., “dry cat food”), we enhance workflow generality by abstracting out example-specific contexts, i.e., replacing “dry cat food” with a more general name “{product-name}” by specifying this in the workflow induction prompts. These workflows are segmented (based on double-line breaks in the model output) and stored separately in the workflow memory. See §A for the model prompts, example

Figure 2: Illustration of our AWM pipeline: an agent takes actions to solve given queries, induces workflows from successful ones, and integrates them into memory.

After the workflows $\mathcal{W}$ are induced, they are then integrated into the agent as auxiliary memory, $M + \mathcal{W} \rightarrow M_w$. M stands for the original agent memory, and M_w stands for the agent memory augmented with induced workflows. When solving a given instruction q , the agent now produces a series of actions by $L(q, M_w, o) = L(q, M + \mathcal{W}, o) \rightarrow a$. In the following, we introduce AWM in use in two scenarios:

SECTION [sec_11]: We also explore a rule-based workflow induction method. See §B for more detailed experiments.

We also explore a rule-based workflow induction method. See §B for more detailed experiments.

Offline Scenario AWM can operate in an *offline* scenario when additional canonical experiences are available, such as data annotated by humans or synthesized by models. In this case, we perform workflow induction and utilization in two standalone processes. As seen in Figure 3, AWM first takes in all training examples from a website by concatenating them into a single prompt, and feeds them to the LM to create a set of workflows at ‘training’ time; $I(\mathcal{E}_{train}) \rightarrow \mathcal{W}_{offline}$. Second, AWM incorporates all induced workflows into the agent memory at inference time to solve test instructions $L(q, M + \mathcal{W}_{offline}, o_i^{test}) \rightarrow a_i^{test}$. Since the workflows are fully induced before test-time inference, the agent uses the same workflow memory $\mathcal{W}_{offline}$ to solve each test.

Online Scenario Extra canonical experiences are not always available or easy to collect, especially those that cover the same domains and tasks as the test instructions. AWM also works in an online, supervision-free setting, where only test queries are needed. Agents with AWM_{online} process test queries in a streaming fashion, where the agents conduct the loop of induce, integrate, and utilize workflows after running inference for each test task.

Concretely, the agent starts with the default memory M . Given the t -th test instruction q_t passed to the agent, the agent attempts to solve the instruction by iteratively forming an experience e_t and predicting a binary label $L_{eval}(e_t) \in \{0, 1\}$. If the prediction is successful, the agent transforms the experience into a workflow w_t and adds it to the agent memory M . This process is repeated until the instruction is solved. The agent then processes the next instruction q_{t+1} . The transcribed equation is AWM_{online} .

3 EXPERIMENTS

3.1 WEBARENA

In this section, we experiment on two major web navigation benchmarks – WebArena (§3.1) and Mind2Web (§3.2). For each benchmark, we first introduce the benchmark and top-performing baseline methods, then present our AWM approach and showcase its ability to achieve superior task success and generalization across varied setups.

For both benchmarks, we conduct AWM on a website basis. In other words, we group examples by their associated websites, and respectively run AWM on each group. This mechanism maintains a small collection of workflows that are nonetheless relevant to the test tasks.

3.1 WEBARENA

WebArena (Zhou et al., 2024) provides 812 web navigation tasks on five websites that cover four common application domains: e-commerce, social forum discussions, collaborative software development, and content management. Most importantly, WebArena supports rigorous evaluation on the functional correctness of agent trajectories.

We adopt the current state-of-the-art method without human-annotated site-specific knowledge, BrowserGym (Drouin et al., 2024), which altered the BrowserGym framework and its default action space, and represent webpages using accessibility trees, following the environment representation in Zhou et al., (2024). Because BrowserGym inputs both webpage HTML and accessibility tree representations, to keep a fair comparison with our method, we also run the BrowserGym version with only accessibility tree webpage representations, denoted as BrowserGym_{ax-tree}. We also compare to the SteP method (Sodhi et al., 2023), which uses 14 human expert written workflows tailored to solving WebArena. Our method, in contrast, uses no human supervision and is not tailored to the WebArena setting.

PARA [para_236]: Table 1: Task success rate on ...

Table 1: Task success rate on WebArena using gpt-4, and score breakdown on five website splits.

Method	Reddit	GitHub	Maps	# Steps
With human engineered workflows				
*SteP (Sodhi et al., 2023)	33.0	37.0	30.0	-
Autonomous agent only				
WebArena (Zhou et al., 2024)	14.9	14.0	11.0	6.0
AutoEval (Pan et al., 2024)	20.2	25.5	18.1	25.4
BrowserGym (Drouin et al., 2024)	23.5	-	-	-
AWM (Ours)	35.5	30.8	29.1	50.9

Following baseline approaches, we use GPT-4 (gpt-4-0613) with a temperature of 0.0 to ensure mostly stable model outputs. Because WebArena only has test examples and no additional high-quality, domain-aligned examples exist, we only conduct AWM in the online setting as in §2.3

3.1.1 MAIN RESULTS


As shown in Table 1, our AWM achieves the best published results on WebArena, surpassing the BrowserGym baseline by 12.0 absolute points and 51.1% relative increase in overall success rate. Notably, AWM also outperforms SteP, which uses strong domain-specific supervision from human-written workflows, by a 7.6% relative increase in overall success rate. According to the breakdown on five websites, our AWM method substantially enhances the agent performance across all websites over the BrowserGym baseline, by 11.8–30.7 absolute points, indicating its general applicability across varied domains and tasks.

Beyond task success, we also evaluate the average number of steps the agent takes to solve a task, as shown in the rightmost column in Table 1. AWM conducts about 2.0 fewer steps per example than the BrowserGym baseline. Further compared to the Autoeval (Pan et al., 2024) method, which necessitates additional evaluation and refinement steps to solve tasks correctly, our AWM approach uses 40.8 fewer steps on average. Both comparisons show that AWM obtains high success rates while maintaining efficient trajectories.

3.1.2 EFFICIENT LEARNING FROM SMALL AMOUNTS OF DATA

To demonstrate the behavior of the AWM_{online} method, we illustrate the cumulative success rate over the process of online evaluation, by evaluating the average success rate of the first k finished examples.

As in Figure 5, the agent exhibits a fast learning curve in the beginning (between 0–40 examples), by acquiring the most essential workflows, which results in higher success rates. Afterward, agents learn more advanced workflows (Figure 1), while success rates gradually stabilize to the highest point in the early learning phase. This showcases AWM’s efficient learning process, which substantially improves performance with merely tens of examples.



3.1.3 CROSS-TEMPLATE WORKFLOW GENERALIZATION

Some tasks in the WebArena benchmark have highly overlapping canonical trajectories, due to the benchmark construction process that instantiates multiple examples from a single underlying task template. AWM intuitively improves in-template success rate, that is, given one workflow induced from a successful example, it would be theoretically easier to solve all other examples generated from the same task template.

PARA [para_282]: Table 2: Task success rate on ...

Table 2: Task success rate on the cross-template subset of WebArena, as well as the result breakdown on each website split. We mark the number of examples for each website split under the name.

Method	Total SR	Shopping (51)	CMS (45)	Reddit (24)	GitLab (45)	Maps (32)
With human engineered workflows						
*ScP (Sodhi et al., 2023)	26.5	20.5	29.3	52.2	27.3	36.4
Autonomous agent only						
AutoEval (Pan et al., 2024)	23.2	12.2	17.1	21.7	31.8	36.4
BrowserGym _{ax-tree}	20.5	10.4	17.8	23.1	27.3	28.6
AWM (OURS)	33.2	24.5	29.3	52.2	31.8	39.4

PARA [para_288]: To confirm that the benefits of ...

To confirm that the benefits of AWM are not just from learning workflows that help only within a template, and investigate whether AWM can obtain cross-template ($\approx$ cross-task) generalization, we extract a subset of WebArena examples sourcing from non-overlapping templates, by grouping examples by their templates and randomly choosing one example from each template group. We run AWM on this cross-template subset and examine if it achieves similar performance gains.

As shown in Table 2, AWM still achieves the highest performance, overall and on each website split. These results demonstrate that AWM induced workflows can effectively generalize across different tasks, i.e., examples instantiated from different task templates.

Building increasingly complex workflows To more intuitively demonstrate AWM’s cross-template generalization and ability to build increasingly complex workflows (as exemplified in Figure 1), we conduct a case study to illustrate the workflow mechanism behind it.




Figure 6: AWM builds increasingly complex workflows over time by learning from past examples and earlier workflows.

As exemplified in Figure 6 the agent creates and learns the “Find a place by its name” workflow in the early stage of the online process by summarizing past examples. Later, when encountering an example that further asks to obtain the zip code of the location, AWM agent learns to adopt the first few steps to find locations by following the existing workflow, and then conducts further steps to obtain the zip code of the place found. Integrating these new steps upon the vanilla find location task, the agent successfully builds a more complex workflow, i.e., “get the zip code of a place”.

3.2 MIND2WEB

Mind2Web (Deng et al., 2023) features web navigation in cross-task, website, and domain settings, stressing the generality of agents on versatile operations and environments. Each task in Mind2Web has a fixed number of steps; at each step, the agent needs to predict an action, which is evaluated by: (1) *element accuracy*: to check if the correct page element is selected, (2) *action F_1* to check if the action taken on the element is correct, and aggregating (1) and (2) yields (3) *step success rate* which checks that both element and action selection are correct at the current step. Lastly, after completing every step in the given task, the last metric (4) *task-level success rate* measures if all intermediate steps are successfully conducted for this task, i.e., all steps for this task score 1 under metric (3).

Because Mind2Web provides a training set covering part of the tested websites (cross-task split), we explore both the *offline* setting that induces workflows from the training set and applies to test sets, and the *online* setting, where we stream workflow induction and inference on test queries (§2.3).

Since we focus on LM-based agents, we compare AWM to two state-of-the-art methods, MindAct (Deng et al., 2023) and Synapse (Zheng et al., 2024). MindAct introduces webpage element filtering and multi-choice task format to ease observation processing, and Synapse changes the format to a trajectory style and augments retrieved relevant examples. We integrate the element filtering adopted in both methods, and added workflows instead of retrieved examples in

PARA [para_331]: Synapse, to verify the superio...

Synapse, to verify the superiority of reusable workflows over concrete examples. To fairly compare with all baseline methods, we run AWM with both gpt-3.5-turbo and gpt-4 models with temperature 0.0. We use the same model for neural workflow induction and agent action generation.

3.2.1 MAIN RESULTS

PARA [para_337]: We first run with AWM of-line us...

PARA [para_333]: Table 3: AWM of-line results on...

We first run with AWM offline using both GPT variants, and find that AWM consistently obtains the highest success rate in both step and task levels, leading to 4.0–8.9% relative and 0.4–2.8 absolute points increases in step-wise and task-wise success rates than the baselines — Synapse with gpt-3.5-turbo and MindAct with gpt-4. Decomposing the step success rate by element and action selection and accuracy, we notice the increases mainly come from more accurate element selection, as indicated by the 5.0

Table 3: AWM offline results on Mind2Web cross-task dataset. Elem Acc and SR are short for element accuracy and success rate. We footnote the GPT variant used by each method, 3.5 and 4 stands for gpt-3.5-turbo and gpt-4, respectively. We highlight the best result within the same model variant.

Method	Elem Acc	Action F ₁	Step SR	SR
MindAct _{3.5}	20.3	56.6	17.4	-
CogAgent _{3.5}	-	-	18.6	0
Synapse _{3.5}	34.0	-	30.6	2.4
AWM _{3.5}	39.0	60.6	36.2	2.8
MindAct ₄	41.6	60.6	36.2	2.0
AWM ₄	50.6	57.3	45.1	4.8

COMMENT [ann_312]:

HIGHLIGHT [ann_309]: (Zheng et al., 2024) me

Abstract sub-routines vs. concrete experiences More specifically, compared to the Synapse (Zheng et al., 2024) method that retrieves the most relevant training examples, AWM achieves a +5.0 element accuracy and leads to a +4.0 increase in step success rate. While augmenting concrete, full examples may bias agents to select elements similar to those presented in the given examples, AWM introduces less bias on element selection via its abstract representation of example-specific

PARA [para_340]: Furthermore, AWM integrates fre... enables higher step success rates.

Furthermore, AWM integrates frequently-used sub-routines, which can be more flexibly and readily leveraged across test examples, compared to the full example trajectories used by Synapse, which are less likely to appear multiple times. In general, our results indicate that the abstract, reusable nature of workflows contributes to the superiority of the AWM method.

HIGHLIGHT [ann_308]: Learn to diverge from workflow guidelines

Learn to diverge from workflow guidelines Despite more accurate element selection, AWM gets slightly lower action F₁ scores than MindAct, possibly because the augmented workflows may guide the agent to take certain actions aligning to the workflows, which are not always relevant to the particular environment state at hand. While following the workflows generally results in more successful task trajectories, agents still encounter some challenges in identifying places to diverge from the workflow guidelines.

PARA [para_342]: 3.2.2 ONLINE AWM ENABLES GENERALIZATION

3.2.2 ONLINE AWM ENABLES GENERALIZATION

Beyond the offline induction setting, we further explore the AWM in the online setting, similar to the WebArena experiment setup in §3.1 where no additional training examples are needed besides test queries. This naturally facilitates cross-website and cross-domain generalization, which we examine on the two other splits provided by the Mind2Web dataset: cross-website and cross-domain tests.

In addition to the MindAct baseline, we additionally set bars with our AWM_{offline} setup, by randomly selecting workflows induced from the training set as memory augmentations. Specifically, for cross-website examples, we select workflows from the same domain; for the cross-domain setting, we randomly select workflows from all domains. We conduct AWM_{online} by iteratively inducing, integrating, and utilizing workflow

COMMENT [ann_315]: SC.

As shown in Table 4 both AWM_{online} and AWM_{offline} surpass the MindAct baseline by a large margin, resulting in 7.4–8.9, 3.6–3.8, and 14.0–16.9 absolute point improvements in step success rates, in cross-task, cross-website, and cross-domain scenarios.

PARA [para_346]: In-domain, cross-task scenario...

In-domain, cross-task scenario When tested in-domain, AWM_{online} and AWM_{offline} perform comparably to each other. When inspecting the model behaviors in detail, we notice the pros and cons of each method. AWM_{online} induces workflows from model-predicted trajectories that are

PARA [para_347]: eqn_306]: The transcribed equation is: ...

PARA [para_362]: Table 4: Success rate on Mind2...

Table 4: Success rate on Mind2Web cross-task, cross-website, and cross-domain generalization test, using gpt-4o model. EA is short for element accuracy and F_1 is short for action F_1 .

Method	Cross-Task				Cross-Website				Cross-Domain				SR
	EA	F_1	Step	SR	EA	F_1	Step	SR	EA	F_1	Step	SR	
AWM _{offline}	50.6	57.3	45.1	4.8	41.4	46.2	33.7	2.3	36.4	41.6	32.6	0.7	
AWM _{online}	50.0	56.4	43.6	4.0	42.1	45.1	33.9	1.6	40.9	46.3	35.5	1.7	

not always correct, thus can lead to incorrect workflows that degrade model performance. On the other hand, the training and test examples on some websites vary in task distributions (e.g., training examples cover how to buy items on Amazon, test examples ask for job applications to Amazon careers.). AWM_{online} naturally resolves this train-test gap because its operating process only involves test queries and environments, therefore yields workflows that are presumably more targeted toward the test distribution, which in turn, leads to higher success rates overall. Nonetheless, if distribution-matching, high-quality training examples are available, AWM_{offline} could bring more benefit by alleviating the gap issue, as the slightly higher cross-tasks scores of AWM_{offline} in Table 4.

Extending to unseen websites and domains When applied on unseen websites or domains, AWM_{online} demonstrates greater generalization abilities, compared to AWM_{offline}. The performance margin of AWM_{online} (over AWM_{offline}) widens as the domain gaps between training and testing data widen from different websites (e.g., *apple* to *bestbuy*) to different domains (e.g., *macys* in shopping domain to *reddit* in social media domain). Because AWM_{online} does not require nor rely on information from the training data, it is not affected by any domain gaps. Nonetheless, as demonstrated by the substantial improvements of AWM_{offline} over the MindAct baseline, AWM_{offline} still demonstrates that models can benefit from mechanistically similar workflows from the previously induced workflow repository.

SECTION 4: HOW MUCH DOES THE SUB-ROUTINE ABSTRACT FORMAT CONTRIBUTE?

4 EXPLORING OPTIMAL WORKFLOW REPRESENTATIONS

In this section, we experiment with other possible alternatives to better represent the workflows. Specifically, we ablate workflows in sub-routine, abstract formats (§4.1), explore workflows in descriptive texts (§4.2), and lastly, beyond the default workflows that describe environment state in NL, we compare strengthened observations with website HTML within workflow steps (§4.3).

4.1 HOW MUCH DOES THE SUB-ROUTINE, ABSTRACT FORMAT CONTRIBUTE?

In this section, we compare our abstract, sub-routine-based induction method using LMs to a rule-based method without context and sub-routine abstraction.

Specifically, our rule-based induction I_{rule} first extracts the action sequence (e.g., *CLICK* → *CLICK* → *TYPE*) of each experience and deduplicates experiences by their action sequence. In each unique experience, we then remove the steps whose action cannot be executed on the environment. We take these unique, validated experiences as workflows. Find more detailed descriptions in §B.

WebArena Results As shown in Table 5 using rule- and LM-based workflow induction performs comparably, with a small 0.1 gap in success rate; the LM-based method appears more efficient and uses 0.4 fewer steps. Our manual analysis found workflows produced by the LM-based induction module I_{lm} are finer-grained, preventing agents from following unnecessary steps that sometimes appear in rule-induced workflows, hence making the task-solving process slightly more efficient.

Table 5: AWM success rate on WebArena using gpt-4o, with rule- and LM-based induction steps

Method	Total SR	# Steps
AWM _{rule}	35.5	6.3
AWM _{lm}	35.5	5.9

Mind2Web Results As shown in Table 6 compared to AWM_{rule}, AWM_{lm} improves by a 2.8 margin. While augmenting concrete, full examples may bias agents to select elements similar to those presented in the given examples, AWM_{lm} introduces less bias on element selection via its abstract representation of example-specific contexts in workflows.

PARA [para_360]: The equation is $\sqrt{(8/1)}$.

PARA [para_398]: Further, AWM [para_400]: uses frequently-

PARA [para_399]: Im
Further, AWM [77] uses frequently-used sub-routines, which can be more flexibly and readily utilized across test examples, compared to the full example trajectories induced by AWM_{rule}, which are less likely to appear multiple times. In general, our results indicate that the abstract, reusable nature of workflows contributes to

SECTION [sec_42]: 4.2 WORKFLOWS IN DESCRIPTIVE TEXTS

4.2 WORKFLOWS IN DESCRIPTIVE TEXTS

AWM represents workflow steps in a program format. In this section, we compare with a textual format for workflows, to understand whether text or code serves as a better format for agent memory. More concretely, we prompt gpt-3.5-turbo to verbalize the action trajectory in the workflows induced in earlier experiments. For example, from an action `CLICK({submit-id})`, its verbalized NL representation reads similar to “CLICK the submit button”. We use the same textual observation and thought from code actions as observation and thoughts in these text actions.

From the results in Table 7, AWM_{text} achieves slightly higher element selection accuracy and step success rate, by 0.6 and 0.3 points, respectively, yet degrades 1.2 in task success rate. Overall, we do not find substantial performance variance between workflows represented in text and code formats, indicating that both forms can bring effective augmentations to agent memory.

SECTION [sec_43]: 4.3 ENVIRONMENT ABSTRACTION IN WORKFLOWS

4.3 ENVIRONMENT ABSTRACTION IN WORKFLOWS

AWM describes intermediate webpage states using NL, yet showing concrete states may be helpful to filter the ground agents on the environment. Since a webpage’s full HTML can be overly long, we filter the webpage representation using the relevance predictor of Deng et al. (2023), and augment each workflow step with this shortened HTML that only has elements predicted as relevant. We run gpt-3.5-turbo with only NL descriptions, only HTML, and both types of content.

As shown in Table 8, NL description of states is more useful than HTML, as replacing NL with HTML leads to a slight 0.8 drop in step success rate. Interestingly, using both NL and filtered HTML leads to worse results. We conjecture the reason to be two-fold. First, adding NL and HTML substantially increases the context length, thus making it harder for models to handle things correctly. Second, the filtered HTML has a substantial number of irrelevant items (missing all correct elements 47% of the time) thus potentially contradicting NL descriptions and impairing agent abilities.

PARA [para_422]: filtered HTML has a substantial... it harder for models to handle things correctly. Second, the filtered HTML has a substantial number of irrelevant items (missing all correct elements 47% of the time) thus potentially contradicting NL descriptions and impairing agent abilities.

HIGHLIGHT [ann_356]: EXPLORING WORKFLOW UTILIZATION IN CONTEXT AND IN ACTION

5 EXPLORING WORKFLOW UTILIZATION IN CONTEXT AND IN ACTION

Besides integrating workflows as agent memory, we also explore workflows in expanding the agent action space, denoted as AWM_{AS}. We leverage the programmatic nature of workflows and wrap each workflow into a high-level function, similar to a shortcut tool the agent can call to perform a pre-determined series of actions (Wang et al. 2024b). Formally, an agent is initially equipped with default, primitive actions P (e.g., `click`, `type`), and AWM_{AS} adds the induced workflow actions W (e.g., `the_agent_calls_get_zipcode`) to its action space.

The agent can call a primitive or workflow action at each step. When a primitive action is called, the agent immediately takes that action. When the agent calls a workflow action, it will trigger the

SECTION [sec_45]: Table 6: AWM results with different workflow induction methods on Mind2Web cross-task dataset.

Table 6: AWM results with different workflow induction methods on Mind2Web cross-task dataset.

Method	Element Acc	Action F1	Step SR	SR
MindAct ₄	41.6	60.6	36.2	2.0
AWM _{4,rule}	49.5	57.0	43.4	2.0
AWM _{4,code}	50.6	57.3	45.1	4.8

PARA [para_402]: the efficacy of AWM [77] method, [77] method, [77] method

PARA [para_403]: the efficacy of AWM [77] method, [77] method, [77] method

Table 7: Mind2Web cross-task results with AWM using code and text workflows.

Method	Element Acc	Action F1	Step SR	SR
AWM _{code}	50.6	57.3	45.1	4.8
AWM _{text}	50.6	57.3	45.1	4.8
AWM _{code+text}	51.2	57.4	45.4	3.6

PARA [para_413]: memory

Table 8: Mind2Web results using GPT-3.5-turbo with different environment representations.

Desc	HTML	Both	Step SR	SR	
✓	✓	✓	39.0	52.8	2.8
✓	✓	✓	38.1	54.0	2.8
✓	✓	✓	37.1	51.3	2.0

PARA [para_421]: it harder for models to handle...

it harder for models to handle things correctly. Second, the filtered HTML has a substantial number of irrelevant items (missing all correct elements 47% of the time) thus potentially contradicting NL descriptions and impairing agent abilities.

COMMENT [ann_364]: Good stuff

Good stuff!

PARA [para_426]: ann_355]: The equation is: $\sqrt{(9/1)}$

PARA [para_448]: sequence of pre-determined steps...

PARA [para_445]: Table 9: Mind2Web results with...

sequence of pre-determined steps in
PARA [para_441]: login (username, password) work-
flow action results in sequentially executing
click(box1-id) → type(box1-id,
username) → click(box2-id)
→ type(box2-id, password) →
click(submit-button). The workflow
action is completed when all intermediate primitive actions are finished.

Table 9: Mind2Web results with AWM_{AS} variant
that alters the action space besides memory augmen-
tation. All methods use gpt-4.

Method	Env Acc	Action F1	Step SR	SR
MindAct	41.6	60.6	36.2	2.0
AWM	50.6	57.3	45.1	4.8
AWM _{AS}	51.8	56.7	46.4	3.6

In Table 9 expanding the agent action space with workflows (AWM_{AS}) slightly improves the step success rate by 1.3 points, and gets the same overall success rate, 3.2, of the base memory-augmented AWM. We analyzed agent predictions and found they call workflow actions in merely 18.5% of the tasks, suggesting a resistance of current agents to use newly-added actions. Overall, expanding actions with workflows seems to reinforce workflows in memory, and brings small extra gains as auxiliary actions.

However, workflow actions do not always lead to task success. A representative example is shown in Figure 7.

When booking flights, users often input a city name such as “New York,” yet the system often pops up some nearby airports to support next-step search. While one can induce a book_flight workflow that enters all required data via a pre-determined action sequence, the action to choose non-up airports is not executable...




Figure 7: An example of dynamic environment changes that challenge workflow action utilization

without seeing the intermediate states with available pop-up options, and is not flexible enough to do so. More advanced techniques such as granting real-time state access or dynamic execution loops can be promising to solve this issue, and we encourage future work to leverage the AWM framework to explore these.

6 RELATED WORK

Web Agent Benchmarks The first modern and widely used web agent benchmark is Shi et al. (2017)’s MiniWeb, which evaluates across various scenarios such as flight booking. (Liu et al. 2018) then created MiniWeb++ with extra challenges. More recently, WebShop (Yao et al. 2022) features a simulated e-commerce website and crowd-sourced text instructions. WebArena (Zhou et al. 2024) integrates four more websites and enables realistic execution-based evaluations, and VisualWebArena (Koh et al. 2024) extends with tasks that necessitate visual inputs. Mind2Web (Deng et al. 2023) proposes versatile tasks and stresses agent generalization across websites and domains. We use WebArena and Mind2Web to evaluate our method’s task success and generality.

Enhancing Agents for Complex Tasks Many works improve agents by modifying their action space, such as constraining its action search space (Liu et al. 2018), enabling LM self-feedback to refine predicted actions (Sun et al. 2023), or incorporating human-designed actions to certain tasks (Sodhi et al. 2023). Other works explore ways to augment agent memory, such as adding example demonstrations in context (Haluptzok et al. 2023, Zheng et al. 2024, Fu et al. 2024). However, high-quality examples are not always available or easy to collect. Our AWM can flexibly operate even when auxiliary examples are non-existent and only test queries are available.

Learning Common Procedures from Experiences Some works use full examples (Zheng et al. 2024) as context for an agent, yet they entangle with example-specific contexts and face challenges in extrapolating to other tasks or domains (Majumder et al. 2023). Many works propose to extract frequently reused sub-routines from experiences with rule-based (Ellis et al. 2023, Bowers et al.

PARA [para_495]: 2023; Grand et al., 2023) or L...

2023; Grand et al., 2023) or LM-based methods (Cai et al., 2023; Wang et al., 2024c[a]) methods, and use them as auxiliary skills to ease future task-solving (Oh et al., 2017; Liang et al., 2023; Yu et al., 2023; Mao et al., 2023). We explored both rule- and LM-based methods to induce reusable workflows, and use them flexibly as context guidance that are free of environment grounding issues.

SECTION [sec_497]: CONCLUSION

7. CONCLUSION

PARA [para_497]: We propose agent work-ow memor...

We propose agent workflow memory that induces, augments, and uses workflows, offline from available examples or purely online at inference time. We evaluate AWM on WebArena and Mind2Web, and achieve 24.6% and 51.1% relative increases in task success rate. AWM also demonstrates its superior generalization abilities across tasks, websites, and domains. We hope AWM sheds insight on and boosts advances in dynamic memory building and agent adaptations on varied digital tasks.

SECTION [sec_499]: ACKNOWLEDGMENTS

ACKNOWLEDGMENTS

PARA [para_499]: We thank Frank Xu, Jiayi Pan, ...

We thank Frank Xu, Jiayi Pan, Vijay Viswanathan, Chenglei Si, and Jason Wu for their helpful discussions during the early stage of this project. Zora Zhiruo Wang is supported by the CMU Teng Family Presidential Fellowship.

REFERENCES

- Matthew Bowers, Theo X. Olausson, Lionel Wong, Gabriel Grand, Joshua B. Tenenbaum, Kevin Ellis, and Armando Solar-Lezama. Top-down synthesis for library learning. *Proc. ACM Program. Lang.*, 7(POPL), jan 2023. doi: 10.1145/3571234. URL <https://doi.org/10.1145/3571234>.
- Tianle Cai, Xuezhi Wang, Tengyu Ma, Xinyun Chen, and Denny Zhou. Large language models as tool makers. *arXiv preprint arXiv:2305.17126*, 2023. URL <https://arxiv.org/pdf/2305.17126>.
- Micheline TH Chi, Paul J Feltovich, and Robert Glaser. Categorization and representation of physics problems by experts and novices. *Cognitive science*, 5(2):121–152, 1981.
- Micheline TH Chi, Robert Glaser, and Marshall J Farr. *The nature of expertise*. Psychology Press, 2014.
- Xiang Deng, Yu Gu, Boyuan Zheng, Shijie Chen, Samuel Stevens, Boshi Wang, Huan Sun, and Yu Su. Mind2web: Towards a generalist agent for the web. In *Thirty-seventh Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2023. URL <https://openreview.net/forum?id=kiYqbO3wqw>.
- Alexandre Drouin, Maxime Gasse, Massimo Caccia, Issam H Laradji, Manuel Del Verme, Tom Marty, Léo Boisvert, Megh Thakkar, Quentin Cappart, David Vazquez, et al. Workarena: How capable are web agents at solving common knowledge work tasks? *arXiv preprint arXiv:2403.07718*, 2024.
- Kevin Ellis, Lionel Wong, Maxwell Nye, Mathias Sable-Meyer, Luc Cary, Lore Anaya Pozo, Luke Hewitt, Armando Solar-Lezama, and Joshua B Tenenbaum. Dreamcoder: growing generalizable, interpretable knowledge with wake–sleep bayesian program learning. *Philosophical Transactions of the Royal Society A*, 381(2251):20220050, 2023.
- Yao Fu, Dong-Ki Kim, Jaekyeom Kim, Sungryull Sohn, Lajanugen Logeswaran, Kyunghoon Bae, and Honglak Lee. Autoguide: Automated generation and selection of state-aware guidelines for large language model agents. *arXiv preprint arXiv:2403.08978*, 2024.
- Gabriel Grand, Lionel Wong, Matthew Bowers, Theo X Olausson, Muxin Liu, Joshua B Tenenbaum, and Jacob Andreas. Lilo: Learning interpretable libraries by compressing and documenting code. *arXiv preprint arXiv:2310.19791*, 2023.
- Patrick Haluptzok, Matthew Bowers, and Adam Tauman Kalai. Language models can teach themselves to program better. In *The Eleventh International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=SaRj2ka1XZ3>.

EQUATION [ann_eqn_416]: The equation is: 11

Summary

COMMENT [ann_432]: Most people correlate pure tool usage with agentic...

COMMENT [ann_433]: agents we need some planning and reasoning capabilities, but IMHO that is only a subset of it. If we want to deploy

COMMENT [ann_434]: other hand comes in many flavors. If we want to go...

COMMENT [ann_435]: After all, an intelligent system should be able to...

After all, an intelligent system should be able to extract and learn from past behavior.

PARA [para_518]: Why do we want to implement a ...

Why do we want to implement a memory module?

- COMMENT [ann_437]: When presented with a new task, a common way to in...
- COMMENT [ann_438]: When presented with a new task, a common way to induce knowledge about the task is to perform in-context learning
- COMMENT [ann_439]: This approach works for a given task but fails to ...
- COMMENT [ann_440]: For example, planning a trip to a foreign country
- COMMENT [ann_441]: differences in the booking procedure
- COMMENT [ann_442]: inspired by how humans abstract common task routin...
- COMMENT [ann_443]: memory module where each workflow represents a goal
- COMMENT [ann_444]: trajectories. Each workflow represents a goal with a common routine extracted from available action

PARA [para_528]: Agent Workflow Memory

Agent Workflow Memory

COMMENT [ann_445]: Memory M, where the base memory contains documenta...

COMMENT [ann_447]: the algorithm looks like

COMMENT [ann_448]: UNTIL

COMMENT [ann_450]: To solve a task specified by an instruction in nat...

- COMMENT [ann_451]: Transition function
- COMMENT [ann_452]: At each time step t_i , the environment state s_i gives...
- COMMENT [ann_453]: generate action a_i via $\pi(q, M, b_i)$ state s_i gives observation o_i . This observation is then passed into the model to
- COMMENT [ann_454]: The action is executed in the environment and chan...
- The action is executed in the environment and changes the state as $T(s_i, a_i) \rightarrow s_{i+1}$

COMMENT [ann_455]: Each completed task forms an experience. This expe...

COMMENT [ann_456]: that we store it includes

- COMMENT [ann_457]: The instruction provided: q
- COMMENT [ann_458]: The workflow contains the steps used to solve th...
- COMMENT [ann_459]: taken $P = (p, a)$
- COMMENT [ann_460]: An induction module is to induce useful workflows ...
- COMMENT [ann_461]: collected examples. These workflows are then ad...


PARA [para_558]: LM-based Workflow Induction

LM-based Workflow Induction

L induces a set of workflows ...

- COMMENT [ann_495]: Module L prompts the agent to extract common subroutines from one or more past agent experiences E .
- COMMENT [ann_496]: If we want to extract common patterns from different routines from one or more input experiences.
- COMMENT [ann_497]: For example, instead of prompting "buy a cover for...", we need to provide more fine-grained instructions.
- COMMENT [ann_498]: For example, instead of prompting "buy a cover for...", the prompt should be "buy a black leather cover for my pixel 9".
- COMMENT [ann_499]: After the workflows W are induced, they are then integrated into the agent as auxiliary memory, $M + W \rightarrow Mw$.
- COMMENT [ann_500]: When solving a given instruction q , the agent then produces a series of actions defined by $L(q, Mw, o) = L(q, M + W, o) \rightarrow a$.

VISUAL: figure: The figure is an illustration ...



PARA [para_567]: Offline Induction


Offline Induction

Assume that you have annotated examples that repre...

- COMMENT [ann_508]: Induction involves two standalone processes.
- COMMENT [ann_509]: Concatenate all training examples into a single pr...
- COMMENT [ann_510]: training time (x) can be done offline.
- COMMENT [ann_511]: Incorporates all induced workflows into the agent ...
- COMMENT [ann_512]: Incorporates all induced workflows into the agent memory at inference time to solve test instructions $L(q, M + W)$.

$q_{test} \rightarrow a_{test}$

VISUAL: figure: The figure illustrates a proce...



PARA [para_589]: Online Induction

Online Induction

COMMENT [ann_547]: addresses the question: What if the agent has not ...

COMMENT [ann_547]: where only test queries are needed. Agents process


COMMENT [ann_548]: loop of inducing, integrating, and utilizing workf

COMMENT [ann_548]: loop of inducing, integrating, and utilizing workflows after running inference for each test task

COMMENT [ann_549]: Assume the default memory is M.

- COMMENT [ann_550]: At test timestep t , instruction q_t is passed to th...
- COMMENT [ann_551]: This instruction paired with the trajectory is the...
- COMMENT [ann_552]: Using LLM as a judge, a binary label (0 or 1) is ass...
- COMMENT [ann_553]: Every time an experience is considered a success...
- COMMENT [ann_554]: into the agent memory $M(t) \leftarrow M(t-1) \cup \{e_t\}$, it is transformed into workflow(s) $\{l_e, t\} \rightarrow \{w, t\}$. $\{w, t\}$ is added into the agent memory $M \leftarrow M \cup \{w, t\} \rightarrow M(t+1)$

VISUAL: figure: The figure illustrates a model...



Jing Yu Koh, Robert Lo, Lawrence Jang, Vikram Duvvur, Ming Chong Lim, Po-Yu Huang, Graham Neubig, Shuyan Zhou, Ruslan Salakhutdinov, and Daniel Fried. Visualwebarena: Evaluating multimodal agents on realistic visual web tasks. In *ICLR 2024 Workshop on Large Language Model (LLM) Agents*, 2024. URL <https://openreview.net/forum?id=RPKxrKTJbJ>.

Jacky Liang, Wenlong Huang, Fei Xia, Peng Xu, Karol Hausman, Brian Ichter, Pete Florence, and Andy Zeng. Code as policies: Language model programs for embodied control. In *2023 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 9493–9500. IEEE, 2023.

Evan Zheran Liu, Kelvin Guu, Panupong Pasupat, and Percy Liang. Reinforcement learning on web interactions using flow-guided exploration. In *International Conference on Learning Representations*, 2018. URL <https://openreview.net/forum?id=ryTp3f-0->.

Bodhisattwa Prasad Majumder, Bhavana Dalvi Mishra, Peter Jansen, Oyvind Tafjord, Niket Tandon, Li Zhang, Chris Callison-Burch, and Peter Clark. Clin: A continually learning language agent for rapid task adaptation and generalization. *arXiv preprint arXiv:2310.10134*, 2023.

Jiayuan Mao, Tomás Lozano-Pérez, Joshua B Tenenbaum, and Leslie Pack Kaelbling. Learning reusable manipulation strategies. In *Conference on Robot Learning*, pp. 1467–1483. PMLR, 2023.

Shikhar Murty, Christopher Manning, Peter Shaw, Mandar Joshi, and Kenton Lee. Bagel: Bootstrapping agents by guiding exploration with language. *arXiv preprint arXiv:2403.08140*, 2024.

Junhyuk Oh, Satinder Singh, Honglak Lee, and Pushmeet Kohli. Zero-shot task generalization with multi-task deep reinforcement learning. In *International Conference on Machine Learning*, pp. 2661–2670. PMLR, 2017.

Jiayi Pan, Yichi Zhang, Nicholas Tomlin, Yifei Zhou, Sergey Levine, and Alane Suhr. Autonomous evaluation and refinement of digital agents. *arXiv preprint arXiv:2404.06474*, 2024.

Christopher Rawles, Alice Li, Daniel Rodriguez, Oriana Riva, and Timothy P Lillicrap. Androidinthewild: A large-scale dataset for android device control. In *Thirty-seventh Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2023. URL <https://openreview.net/forum?id=j4b3l5kO1l>.

Christopher Rawles, Sarah Clinckemaulle, Yifan Chang, Jonathan Waltz, Gabrielle Lau, Marybeth Fair, Alice Li, William Bishop, Wei Li, Folawiyo Campbell-Ajala, et al. Androidworld: A dynamic benchmarking environment for autonomous agents. *arXiv preprint arXiv:2405.14573*, 2024.

Tianlin Shi, Andrej Karpathy, Linxi Fan, Jonathan Hernandez, and Percy Liang. World of bits: An open-domain platform for web-based agents. In Doina Precup and Yee Whye Teh (eds.), *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pp. 3135–3144. PMLR, 06–11 Aug 2017. URL <https://proceedings.mlr.press/v70/shi17a.html>.

Paloma Sodhi, SRK Branavan, and Ryan McDonald. Heap: Hierarchical policies for web actions using llms. *arXiv preprint arXiv:2310.03720*, 2023.

Haotian Sun, Yuchen Zhuang, Lingkai Kong, Bo Dai, and Chao Zhang. Adaplaner: Adaptive planning from feedback with language models. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023. URL <https://openreview.net/forum?id=rnKgbKmelt>.

Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *Transactions on Machine Learning Research*, 2024a. ISSN 2835-8856. URL <https://openreview.net/forum?id=ehfRiF0R3a>.

Zhiruo Wang, Zhoujun Cheng, Hao Zhu, Daniel Fried, and Graham Neubig. What are tools anyway? a survey from the language model perspective. In *First Conference on Language Modeling*, 2024b. URL <https://openreview.net/forum?id=Xh1B90iBSR>.

Zhiruo Wang, Graham Neubig, and Daniel Fried. TroVE: Inducing verifiable and efficient toolboxes for solving programmatic tasks. In *Forty-first International Conference on Machine Learning*, 2024c. URL <https://openreview.net/forum?id=DCNCwaMJjI>.

Shunyu Yao, Howard Chen, John Yang, and Karthik Narasimhan. Webshop: Towards scalable real-world web interaction with grounded language agents. In S. Koyejo, S. Mohamed, A. Agarwal, D. Belgrave, K. Cho, and A. Oh (eds.), *Advances in Neural Information Processing Systems*, volume 35, pp. 20744–20757. Curran Associates, Inc., 2022. URL https://proceedings.neurips.cc/paper_files/paper/2022/file/82ad13ec01f9fe44c01cb91814fd7b8c-Paper-Conference.pdf.

Ori Yoran, Samuel Joseph Amouyal, Chaitanya Malaviya, Ben Bogin, Ofir Press, and Jonathan Berant. Assistantbench: Can web agents solve realistic and time-consuming tasks? *arXiv preprint arXiv:2407.15711*, 2024.

Wenhao Yu, Nimrod Gileadi, Chuyuan Fu, Sean Kirmani, Kuang-Huei Lee, Montse Gonzalez Arenas, Hao-Tien Lewis Chiang, Tom Erez, Leonard Hasenclever, Jan Humplik, et al. Language to rewards for robotic skill synthesis. *arXiv preprint arXiv:2306.08647*, 2023.

Longtao Zheng, Rundong Wang, Xinrun Wang, and Bo An. Synapse: Trajectory-as-exemplar prompting with memory for computer control. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=Pc8AU1aF5e>.

Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng, Tianyue Ou, Yonatan Bisk, Daniel Fried, Uri Alon, and Graham Neubig. Webarena: A realistic web environment for building autonomous agents. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=oKn9c6ytLx>.

SECTION [para_664] A. LM-BASED WORKFLOW INDUCTION

A LM-BASED WORKFLOW INDUCTION

COMMENT [ann_610]: S2.3,

As introduced in §2.3, one realization of our workflow induction module is to prompt LMs to generate abstract, sub-routine workflows from the given examples, i.e., experience. In this section, we provide the detailed model prompt, exemplar workflows induced by models, and quality examination on these workflows.

SECTION [para_664] A.1. MODEL PROMPT

A.1. MODEL PROMPT

PARA [para_664]: We provide the exact prompt in...

We provide the exact prompt inputted to the model for WebArena and Mind2Web experiments below. Experiments on both datasets use the same prompt.

PARA [para_665]: Given a list of web navigation...

Given a list of web navigation tasks, your task is to extract the common workflows.

PARA [para_667]: You need to find the repetitive...

Each given task contains a natural language instruction, and a series of actions to solve the task.

You need to find the repetitive subset of actions across multiple tasks, and extract each of them out as a workflow.

Each workflow should be a commonly reused sub-routine of the tasks. Do not generate similar or overlapping workflows. Each workflow should have at least two steps. Represent the non-fixed elements (input text, button strings) with descriptive variable names as shown in the example.

PARA [para_666]: the task.

SECTION [para_668] A.2. EXEMPLAR WORKFLOWS

A.2. EXEMPLAR WORKFLOWS

PARA [para_669]: We present several exemplar wo...

We present several exemplar workflows induced on WebArena and Mind2Web, to give a more concrete impression of workflows.

PARA [para_670]: WebArena Work-ows We show one ...

WebArena Workflows. We show one example workflow on each website involved in WebArena.

PARA [para_671]: ## Shopping: Browse Products...

Shopping: Browse Products in a Specific Category

To browse products in a specific category, I need to navigate to the relevant main category. I will start by hovering over the main category menu item to reveal the subcategories.

PARA [para_673]: hover('main category id'). To b...

hover('main category id')

To browse products in the specific subcategory, I need to click on the subcategory link.

PARA [para_675]: click('subcategory id'). I need to click on the subcategory link...

click('subcategory id')

PARA [para_676]: ## shopping admin: Edit and Sa...

shopping admin: Edit and Save Changes

PARA [para_677]: To edit a specific field, I need...

PARA [para_679]: clear('field id'). To edit a specific field, I need to locate the field and update its value.

PARA [para_680]: fill('field id', 'new value') ...

clear('field id')

fill('field id', 'new value')

PARA [para_682]: click('save button id'). Next, I need to save the changes by clicking the "Save" button.

click('save button id')

PARA [para_683]: ## reddit: Navigate to a forum...

PARA [para_684]: To navigate to a specific forum...

reddit: Navigate to a forum section and select a specific forum

PARA [para_685]: click('.42'). Now, I need to cli...

click('.42')

PARA [para_687]: click('<forum link id>'). Now, I need to click on the specific forum link based on the forum name provided.

click('<forum link id>')

PARA [para_688]: ## gitlab: Navigation to Repos...

PARA [para_689]: To its contributors to find detailed...

gitlab: Navigation to Repository and Contributors Section

PARA [para_690]: contribution data.

This workflow involves searching for a repository and navigating to its contributors to find detailed contribution data.

PARA [para_691]: First, search for the speci-c ...

PARA [para_692]: fill('130', '{RepositoryName}'). First, search for the specific repository to gather information.

fill('130', '{RepositoryName}')

press('130', 'Enter')

PARA [para_694]: click('311'). # Contributors ...

Navigate to the "Contributors" section to view contribution details.

click('311') # "Contributors" link

PARA [para_695]: send_msg_to_user('{Contributor... Obtain and report the required contributor details.

send_msg_to_user('{ContributorDetails}')

PARA [para_696]: The equation is: $\sqrt{\frac{1}{e}}$...

PARA [para_710]: ## map: Calculate Travel Time ...

map: Calculate Travel Time and Distance

To calculate travel time and distance between two locations, I will use the directions feature. I will fill in the respective fields and select the mode of transportation.

PARA [para_712]: fill(158, FROM LOCATION)

PARA [para_713]: fill(163, TO LOCATION)

PARA [para_714]: select option(169, MODE OF

PARA [para_715]: click(171, I will use these

click('171')

PARA [para_717]: send msg to user(The distance

I will use these details to provide the user with accurate travel time and distance information.

PARA [para_718]: TO LOCATION IS DISTANCE and th...

send msg to user(The distance between FROM LOCATION and

TO LOCATION is DISTANCE and the estimated travel time is TIME.)

PARA [para_719]: Mind2Web Workflows We present o...

Mind2Web Workflows. We present one example workflow in each data domain in Mind2Web.

PARA [para_720]: # travel: enter flight locations...

travel: enter flight locations

PARA [para_722]: city/airport for your flight.

Given that you are on the flight

city/airport for your flight.

PARA [para_721]: booking page, this workflow ent...

booking page, this workflow enters the departure and destination

[link] From Departure Airport or City Your Origin - > CLICK

[textbox] Origin City or Airport - > TYPE: {your-origin-city}

[link] {best-popup-option} - > CLICK

[link] To Destination Airport or City Your Destination - > CLICK

[textbox] Destination City or Airport - > TYPE: {your-destination-city}

[link] {best-popup-option} - > CLICK

PARA [para_724]: # shopping: search and sort Gi...

shopping: search and sort

PARA [para_725]: search results.

Given that you are on the Amazon

search results.

PARA [para_727]: [textbox] Search Amazon -> TYP...

sorts the results.

[textbox] Search Amazon - > TYPE: {search-term}

[button] Go - > CLICK

[span] Sort by: - > CLICK

[option] {sort-option} - > CLICK

PARA [para_728]: # entertainment: search and se...

entertainment: search and select

PARA [para_730]: match.

Given that you are on the IMDb homepage,

match.

PARA [para_731]: [textbox] Search IMDb -> TYPE:...

[textbox] Search IMDb - > TYPE: {search-term}

[button] Submit Search - > CLICK

[button] {best-match} - > CLICK

PARA [para_729]: this workflow searches for a te...

PARA [para_726]: this workflow searches for a product and

PARA [para_723]: this workflow searches for a product and

PARA [para_729]: this workflow searches for a te...

PARA [para_726]: this workflow searches for a product and

PARA [para_723]: this workflow searches for a product and

PARA [para_729]: this workflow searches for a te...

PARA [para_726]: this workflow searches for a product and

PARA [para_723]: this workflow searches for a product and

PARA [para_729]: this workflow searches for a te...

PARA [para_726]: this workflow searches for a product and

PARA [para_723]: this workflow searches for a product and

PARA [para_729]: this workflow searches for a te...

PARA [para_726]: this workflow searches for a product and

PARA [para_723]: this workflow searches for a product and

PARA [para_729]: this workflow searches for a te...

PARA [para_726]: this workflow searches for a product and

PARA [para_723]: this workflow searches for a product and

PARA [para_729]: this workflow searches for a te...

PARA [para_726]: this workflow searches for a product and

PARA [para_723]: this workflow searches for a product and

PARA [para_729]: this workflow searches for a te...

PARA [para_726]: this workflow searches for a product and

PARA [para_723]: this workflow searches for a product and

PARA [para_729]: this workflow searches for a te...

PARA [para_726]: this workflow searches for a product and

PARA [para_723]: this workflow searches for a product and

PARA [para_729]: this workflow searches for a te...

PARA [para_726]: this workflow searches for a product and

PARA [para_723]: this workflow searches for a product and

PARA [para_729]: this workflow searches for a te...

PARA [para_726]: this workflow searches for a product and

As shown in Table 10, neural-based induction produces 7.3–7.4 workflows per example, which is efficient and does not add too much content to the memory. On WebArena, the induced workflows are

PARA [para_766]: the equation is: 15

PARA [para_750]: Table 10: Quality evaluation o...

Table 10: Quality evaluation of model-induced workflows on Mind2Web dataset.

Method	WebArena	WebArena	Function Overlap	Utility Rate
WebArena	7.4	-	0.08	0.94
Mind2Web	7.3	0.40	0.20	0.91

PARA [para_753]: used by 0.94 of the test exam...

PARA [para_754]: 0.08 of the steps between work...
 used by 0.94 of the test examples, indicating its wide applicability among varied tasks. Further, only 0.08 of the steps between workflows overlap, demonstrating the efficiency of workflows in solving respective tasks. Workflows on Mind2Web, although used similarly frequently as indicated by the high 0.91 utility rate, have slightly more functional overlap, and only achieve a 0.40 coverage over test examples. However, as the training examples used to induce workflows have substantial task distribution variances with the cross-task test examples, this relatively low coverage is reasonable.

SECTION [sec_52]: RULE-BASED WORKFLOW INDUCTION

B RULE-BASED WORKFLOW INDUCTION

PARA [para_750]: Beyond LM-based workfow induc...

Beyond LM-based workflow induction, we also explored a rule-based workflow induction method. Our rule-based workflow induction module consists of two steps: (i) experience deduplication, and (2) invalid action filtering.

PARA [para_758]: CLICK-TYPE from the trajectory...
 For deduplication, we extract the action sequence of the experience, e.g., extracting CLICK → CLICK → TYPE from the trajectory CLICK('12') → CLICK('30') → TYPE('44', "cat"). We group experiences by their action sequence and randomly select n ($n = 1$ by default) experiences from each group. Specifically on WebArena, where the task template for each experience is available. We conduct another round of deduplication by grouping experiences by their task template, and randomly selecting n ($n = 1$ by default) experiences from each group. This process yielded diverse experiences from the given set of experiences.

Next, for each unique experience, we remove the invalid steps in its action trajectory. Invalid actions means actions that cannot be successfully executed on the environment, because the input arguments do not meet the requirement of the action function. Specifically, we have one rule of determining invalid actions for CLICK and TYPE, that requires the first argument to be a string-formatted integer (which refers to the id of an element in the environment). We remove CLICK and TYPE steps if they do not meet this requirement. For example, an experience with trajectory CLICK(12) → CLICK('12') → CLICK('30') → TYPE(44, "cat") → TYPE('44', "cat") will yield CLICK('12') → CLICK('30') → TYPE('44', "cat"). We conduct this invalid action filtering for each unique experience, and take the resulting experiences as rule-based workflows.

SECTION [sec_53]: INTEGRATING AWM OFFLINE AND ONLINE

C INTEGRATING AWM OFFLINE AND ONLINE

COMMENT [ann_640]: §3.2,

We compared $AWM_{offline}$ and AWM_{online} in §3.2 that adopts workflows induced separately from training or on-the-fly during testing, respectively. In this section, we explore an integration of both sets of workflows, AWM_{off+on} , that injects relevant training workflows to warm start task-solving, but also aggregates increasingly more online-induced workflows to better adapt to test distributions.

PARA [para_763]: Table 11: Success rate on Mind...

Table 11: Success rate on Mind2Web cross-task, cross-website, and cross-domain generalization testing using $AWM_{offline}$, AWM_{online} , and AWM_{off+on} . EA is short for element accuracy, and AF_1 is short for action F_1 .

Method	Cross-Task				Cross-Website				Cross-Domain			
	EA	AF ₁	Step SR	SR	EA	AF ₁	Step SR	SR	EA	AF ₁	Step SR	SR
$AWM_{offline}$	50.6	57.3	45.1	4.8	41.4	46.2	33.7	2.3	36.4	41.6	32.6	0.7
AWM_{online}	50.0	56.4	43.6	4.0	42.1	45.1	33.9	1.6	40.9	46.3	35.5	1.7
AWM_{off+on}	50.0	57.0	44.5	1.6	41.8	45.5	33.3	1.1	39.3	44.3	34.1	1.5

COMMENT [ann_641]: Table 11,

From Table 11 AWM_{off+on} scores between $AWM_{offline}$ and AWM_{online} across three test splits. Rather than an additive effect, workflows induced offline and online are not fully compatible with each other, particularly, the offline workflows seem to impair the generative quality and utility efficacy of online workflows, therefore resulting in medium results overall.

PARA [para_769]: The equation is: $\sqrt{16}$